

textNet: Directed, Multiplex, Multimodal Event Network Extraction from Textual Data

Elise Zufall¹ and ¹

¹ UC Davis, USA

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).

Introduction

A number of tools exist to generate networks based on co-occurrence of words within documents (such as the [Nocodefunctions](#) app ([Levallois et al., 2012](#)), the “[textnets](#)” package ([Bail, 2024](#)), [InfraNodus](#) ([Paranyushkin, 2018](#)), and many more). ## Statement of Need

Directed Graph Production

Multiplex Graph Output

Multimodal Graph Output

Existing packages such as the [manynet](#) package ([Hollway, 2024](#)) the default entity type tags for an NLP engine such as [spaCy](#) ([Honnibal et al., 2021](#)),

Avoids Saturation

Installation

The stable version of this package can be installed from Github, using the [devtools](#) package ([Wickham et al., 2022](#)):

such as [spacyr](#) ([Benoit et al., 2023](#)), [pdftools](#) ([Ooms, 2024](#)), [igraph](#) ([Csárdi et al., 2024](#)), and [network](#) ([Butts et al., 2023](#)). To use the full functionality of [textNet](#), such as pre-processing tools and post-processing analysis tools, we recommend installing these packages, which for [spacyr](#) requires integration with Python. However, the user may wish to preprocess and parse data using their own NLP engine, and skip directly to the [textnet_extract\(\)](#) function, which does not depend on any of the aforementioned packages. The [textnet_extract\(\)](#) function does, however, use functions from [pbapply](#) ([Solymos et al., 2023](#)), [data_table](#) ([Barrett et al., 2024](#)), [dplyr](#) ([Wickham et al., 2023](#)), and [tidyr](#) ([Wickham et al., 2024](#)).

Overview and Main Functions

- [OPTIONAL] Pre-processing: [pdf_clean\(\)](#), a wrapper for the [pdftools::pdf_text\(\)](#) function which includes a custom header/footer text removal feature; and [parse_text\(\)](#), which is a wrapper for the [spacyr](#) package and uses the [spaCy](#) natural language processing engine ([Honnibal et al., 2021](#)) to parse text and perform part of speech tagging, dependency parsing, and named entity recognition (NER). Alternatively, as described below, the user can skip this step and load parsed text directly into the package.
- Network extraction: [textnet_extract\(\)](#), which generates a graph database from parsed text based upon tags and dependency relations
- Disambiguation: tools for cleaning, recoding, and aggregating node and edge attributes, such as the [find_acronyms\(\)](#) function, which can be paired with the [disambiguation\(\)](#) function to identify acronyms in the text and replace them with the full entity name.

- Exploration: the `export_to_network()` function for exporting the graph database to `igraph` and network objects, `top_features()` for viewing node and edge attributes, and `combine_networks()` for aggregating multiple document-based graphs based on common nodes.

Example

Exogenous metadata that has been collected separately by the researcher regarding the different documents and their real-world context. The extracted networks, with their collections of verb attributes, node attributes, edge incidences, and edge attributes, can also be analyzed through a variety of tools, such as an Exponential Random Graph Model, to determine the probability of edge formation under certain conditions. A Temporal Exponential Random Graph Model could also shed light on the changes of a document over time, such as the multiple versions of the groundwater sustainability plan in this example.

Entity Network Extraction Algorithm

The directed network generated by *textNet* represents the collection of all identified entities in the document, joined by edges signifying the verbs that connect them. The user can specify which entity categories should be preserved. The output format is a list containing four data.tables: an edgelist, a nodelist, a verblist, and an appositive list.

The edgelist includes edge attributes such as verb tense, any auxiliary verbs in the verb phrase, whether an open clausal complement (Universal Dependencies code “xcomp”) is associated with the primary verb, whether any hedging words were detected in the sentence, and whether any negations were detected in the sentence.

The returned edgelist by default contains both complete and incomplete edges. A complete edge includes a source, verb, and target. An incomplete edge includes either a source or a target, but not both, along with its associated verb. Incomplete edges convey information about which entities are commonly associated with different verbs, even though they do not reveal information about which other entities they are linked to in the network. These incomplete edges can be filtered out when converting the output into a network object, such as through the *network* package or the *igraph* package. The nodelist returns all entities of the desired types found in the document, regardless of whether they were found in the edgelist. Thus, the nodelist allows the presence of isolates to be documented, as well as preserving node attributes. The verblist includes all of the verbs found in the document, along with verb attributes imported from *VerbNet* (Kipper-Schuler, 2006). This can be used to conduct analyses of certain verb classifications of interest. Finally, the appositive list is a table of entities that may be synonyms. This list is generated from entities whose universal dependency parsing labels as appositives, and whose head token points to another entity. These pairs are included in the table as potential synonyms. If this feature is used, cleaning and filtering by hand is recommended, as appositives can at times be misidentified by existing NLP tools. An automated alternative we recommend is our `find_acronym` tool, which scans the entire document for acronyms defined parenthetically in-text and compiles them in a table.

This network is directed such that the entities that form the subject of the sentence are denoted as the “source” nodes, and the remaining entities are denoted as the “target” nodes. To identify whether each entity is a “source” or a “target”, we use dependency parsing in the Universal Dependencies format, in which each token in a given sentence has an associated “syntactic head” token from which it is derived. Starting with each entity in the sentence, the chain of syntactic head tokens is traced back until either a subject or a verb is reached. If it reaches a subject first, the entity is considered a “source.” If it reaches a verb first, it is considered a “target.”

To identify the subject, we search for the presence of at least one of the following subject tags: “nsubj” (nominal subject), “nsubjpass” (nominal subject – passive), “csubj” (clausal subject),

“csubjpass” (clausal subject – passive), “agent”, and “expl” (expletive). To identify the object, we search for the presence of at least one of the following: “pobj” (object of preposition), “iobj” (indirect object), “dative”, “attr” (attribute), “dobj” (direct object), “oprd” (object predicate), “ccomp” (clausal complement), “xcomp” (open clausal complement), “acomp” (adjectival complement), or “pcomp” (complement of preposition).

If a subject token is reached first (“nsubj”, “nsubjpass”, “csubj”, “csubjpass”, “agent”, or “expl”), this indicates that the original token is doing the verb action. That is, it serves some function related to the subject of the sentence. We designate this by tagging it “source,” since these types of relationships will be used to designate the “from” or “source” nodes in our directed network. If a verb token is reached first (“VERB” or “AUX”), this indicates that the verb action is occurring for or towards the original token, which we denote with the tag “target.” These tokens are potential “to” or “target” nodes in our directed network. Linking the two nodes is an edge representing the verb that connects them in the sentence.

Due to the presence of tables, lists, or other anomalies in the original document, it is possible that a supposed “sentence” has a head token trail that does not lead to a verb as is normatively the case. In these instances, the tokens whose trails terminate with a non-subject, non-verb token are assigned neither “source” nor “target” tags. Finally, an exception is made if an appositive token is reached first, since this indicates that the token in question is merely a synonym or restatement of an entity that is already described elsewhere in the sentence and, accordingly, should not be treated as a separate node. Tokens that lead to appositives are assigned neither “source” nor “target” tags, but are preserved as a separate appositive list.

If a verb phrase in the edgelist does not have any sources, the sources associated with the head token of the verb phrase’s main verb (that is, the verb phrase’s parent verb) are adopted as sources of that verb phrase. As of Version 1.0, *textNet* does not do this recursively, to preserve performance optimization.

The `textNet::textnet_extract()` function returns the full list of open clausal complement lemmas associated with the main verb as an edge attribute: “xcomp_verb”. The list of auxiliary verbs and their corresponding lemmas associated with the main verb, as well as the list of auxiliary verbs and corresponding lemmas associated with the open clausal complements linked to the main verb, are also included as edge attributes: “helper_token”, “helper_lemma”, “xcomp_helper_token”, and “xcomp_helper_lemma”, respectively.

The extraction function also detects hedging words and negations. The function `textNet::textnet_extract()` produces an edge attribute “has_hedge”, which is T if there is a hedging auxiliary verb (“may”, “might”, “can”, “could”) or main verb (“seem”, “appear”, “suggest”, “tend”, “assume”, “indicate”, “estimate”, “doubt”, “believe”) in the verb phrase.

Tense is also detected. The six tenses tagged by *spaCy* in `textNet::parse_text()` are preserved by `textNet::textnet_extract()` as an edge attribute “head_verb_tense”. This attribute can take on one of six values: “VB” (verb, base form), “VBD” (verb, past tense), “VBG” (verb, gerund or present participle), “VBN” (verb, past participle), “VBP” (verb, non-3rd person singular present), or “VBZ” (verb, 3rd person singular present). Additionally, an edge attribute “is_future” is generated by `textNet::textnet_extract()`, which is T if the verb phrase contains an xcomp, has the token “going” as a head verb, and a being verb token as an auxiliary verb (i.e. is of the form “going to ”) or contains one of the following auxiliary verbs: “shall”, “will”, “wo”, or “ll” (i.e. is of the form “will ”).

Acknowledgements

The authors gratefully acknowledge the support of the Sustainable Agricultural Systems program, project award no. 2021-68012-35914, from the U.S. Department of Agriculture’s National Institute of Food and Agriculture and the National Science Foundation’s Dynamics of Integrated Socio-Environmental Systems program, grant no. 2205239.

137 Appendix

138 This appendix describes the pre-processing tools available through the *textNet* package, which
139 enable the user to generate the data frame expected by the `textnet_extract()` function.

140 Pre-Processing Step I: Process PDFs

141 This is a wrapper for `pdftools`, which has the option of using `pdf_text` or `OCR`. We have also
142 added an optional header/footer removal tool. This optional tool is solely based on carriage
143 returns in the first or last few lines of the document, so may inadvertently remove portions
144 of paragraphs. However, not removing headers or footers can lead to improper inclusion of
145 header and footer material in sentences, artificially inflating the presence of nodes whose entity
146 names are included in the header and footer. Because of the risk of headers and footers to
147 preferentially inflate the presence of a few nodes, the header/footer remover is included by
148 default. It can be turned off if the user has a preferred header/footer removal tool to use
149 instead, or if the input documents lack headers and footers.

```
150 library(textNet)
151 library(stringr)
152 URL <- "https://sgma.water.ca.gov/portal/service/gspdocument/download/2840"
153 download.file(URL, destfile = "old.pdf", method="curl")
154
155 URL <- "https://sgma.water.ca.gov/portal/service/gspdocument/download/9625"
156 download.file(URL, destfile = "new.pdf", method="curl")
157
158 pdfs <- c("old.pdf",
159           "new.pdf")
160
161 old_new_text <- textNet::pdf_clean(pdfs, keep_pages=T, ocr=F, maxchar=10000,
162                                   export_paths=NULL, return_to_memory=T, suppressWarn = F, auto_headf
163 names(old_new_text) <- c("old", "new")
```

164 Pre-Processing Step II: Parse Text

165 This is a wrapper for the pre-trained multipurpose NLP model *spaCy* (Honribal et al., 2021),
166 which we access through the R package *spacyr* (Benoit et al., 2023). It produces a table that
167 can be fed into the `textnet_extract` function in the following step. To initialize the session, the
168 user must define the "RETICULATE_PYTHON" path, abbreviated as "ret_path" in *textNet*,
169 as demonstrated in the example below. The page contents processed in the Step 1 must
170 now be specified in vector form in the "pages" argument. To determine which file each page
171 belongs to, the user must specify the `file_ids` of each page. We have demonstrated how to do
172 this below. The package by default does not preserve hyphenated terms, but rather treats
173 them as separate tokens. This can be adjusted.

174 The user may also specify "phrases_to_concatenate", an argument representing a set of
175 phrases for *spaCy* to keep together during its parsing. The example below demonstrates
176 how to use this feature to supplement the NER capabilities of *spaCy* with a custom list
177 of entities. This supplementation could be used to ensure that specific known entities are
178 recognized; for instance, *spaCy* might not detect that a consulting firm such as "Schmidt
179 and Associates" is one entity rather than two. Conversely, this capability could be leveraged
180 to create a new category of entities to detect, that a pretrained model is not designed to
181 specifically recognize. For instance, to create a public health network, one might include a
182 known list of contaminants and diseases and designate custom entity type tags for them, such
183 as "CONTAM" and "DISEASE"). In this example, we investigate the connections between the
184 organizations, people, and geopolitical entities discussed in the plan with the flow of water in
185 the basin. To assist with this, we have input a custom list of known water bodies in the region
186 governed by our test document and have given it the entity designation "WATER". This is

carried out by setting the variable “phrases_to_concatenate” to a character vector, including all of the custom entities. Then, the entity type can be set to the desired category. Note that this function is case-sensitive.

```
library(findpython)
ret_path <- find_python_cmd(required_modules = c('spacy', 'en_core_web_lg'))

water_bodies <- c("surface water", "Surface water", "groundwater", "Groundwater", "Sa
Chowchilla canal", "lower aquifer", "upper aquifer", "upper and lower aquifers", "lower

old_new_parsed <- textNet::parse_text(ret_path,
                                     keep_hyph_together = F,
                                     phrases_to_concatenate = water_bodies,
                                     concatenator = "_",
                                     text_list = old_new_text,
                                     parsed_filenames=c("old_parsed","new_parsed"),
                                     overwrite = T,
                                     custom_entities = list(WATER = water_bodies))
```

Another NLP tool may be used instead of the built-in *textNet* function at this phase, as long as the output conforms to spaCy tagging standards: Universal Dependencies tags for the “pos” part-of-speech column (Nivre, 2017), and Penn Treebank tags for the “tags” column (Marcus et al., 1999). The *textnet_extract* function expects the parsed table to follow specific conventions. First, a row must be included for each token. The column names expected by *textnet_extract* are:

- doc_id, a unique ID for each page
- sentence_id, a unique ID for each sentence
- token_id, a unique ID for each token
- token, the token, generally a word, represented as a string
- lemma, the canonical or dictionary form of the token
- pos, a code referring to the token’s part of speech, defined according to [Universal Dependencies](#) (Nivre, 2017).
- tag, a code referring to the token’s part of speech, according to [Penn Treebank](#) (Marcus et al., 1999).
- head_token_id, a numeric ID referring to the token_id of the head token of the current row’s token
- dep_rel, the dependency label according to [ClearNLP Dependency](#) labels (Choi, 2024)
- entity, the entity type category defined by [OntoNotes 5.0](#) (Weischedel et al., 2012). This is represented as a string, ending in “_B” if it is the first token in the entity or “_I” otherwise).

References

- Bail, C. (2024). *Cbail/textnets* (Version 0.1.1). <https://github.com/cbail/textnets>
- Barrett, T., Dowle, M., Srinivasan, A., Gorecki, J., Chirico, M., Hocking, T., Schwendinger, B., Stetsenko, P., Short, T., Lianoglou, S., Antonyan, E., Bonsch, M., Parsonage, H., Ritchie, S., Ren, K., Tan, X., Saporta, R., Seiskari, O., Dong, X., ... Krylov, I. (2024). *Data.table: Extension of 'data.frame'* (Version 1.16.2). <https://cran.r-project.org/web/packages/data.table/index.html>
- Benoit, K., Matsuo, A., Gruber, J., & Council (ERC-2011-StG 283794-QUANTESS), E. R. (2023). *Spacyr: Wrapper to the 'spaCy' 'NLP' library* (Version 1.3.0). <https://cran.r-project.org/web/packages/spacyr/index.html>
- Butts, C. T., Hunter, D., Handcock, M., Bender-deMoll, S., Horner, J., Wang, L., Krivitsky, P.

- 237 N., Knapp, B., Bojanowski, M., & Klumb, C. (2023). *Network: Classes for relational data*
 238 (Version 1.18.2). <https://cran.r-project.org/web/packages/network/index.html>
- 239 Choi, J. (2024, August 7). *ClearNLP dependency labels*. *ClearNLP guidelines*. https://github.com/clir/clearnlp-guidelines/blob/master/md/specifications/dependency_labels.md
 240
- 241 Csárdi, G., Nepusz, T., Traag, V., Horvát, S., Zanini, F., Noom, D., Müller, K., Salmon, M.,
 242 Antonov, M., & details, C. Z. I. *igraph: Network analysis and*
 243 *visualization* (Version 2.1.1). <https://cran.r-project.org/web/packages/igraph/index.html>
- 244 Hollway, J. (2024). *Manynet: Many ways to make, modify, map, mark, and measure myriad*
 245 *networks* (Version 1.2.6). <https://CRAN.R-project.org/package=manynet>
- 246 Honnibal, M., Montani, I., Van Landeghem, S., & Boyd, A. (2021). *spaCy: Industrial-strength*
 247 *natural language processing in python* (Version 3.1.3). [https://github.com/explosion/](https://github.com/explosion/spaCy/tree/master)
 248 [spaCy/tree/master](https://github.com/explosion/spaCy/tree/master)
- 249 Kipper-Schuler, K. (2006). *VerbNet* (Version 3.3). University of Colorado Boulder. <https://verbs.colorado.edu/verb-index/vn3.3/index.php>
 250
- 251 Levallois, C., Clithero, J. A., Wouters, P., Smidts, A., & Huettel, S. A. (2012). Translating
 252 upwards: Linking the neural and social sciences via neuroeconomics. *Nature Reviews*
 253 *Neuroscience*, 13(11), 789–797. [https://nocodfunctions.com/cowo/semantic_networks_](https://nocodfunctions.com/cowo/semantic_networks_tool.html)
 254 [tool.html](https://nocodfunctions.com/cowo/semantic_networks_tool.html)
- 255 Marcus, M., Santorini, B., Marcinkiewicz, M. A., & Taylor, A. (1999). *Treebank-3 documenta-*
 256 *tion*. <https://doi.org/https://doi.org/10.35111/gq1x-j780>
- 257 Nivre, J. (2017). *Universal POS tags* (Version 2). Universal Dependencies. [https://](https://universaldependencies.org/u/pos/)
 258 universaldependencies.org/u/pos/
- 259 Ooms, J. (2024). *Pdftools: Text extraction, rendering and converting of PDF documents*
 260 (Version 3.4.1). <https://cran.r-project.org/web/packages/pdftools/index.html>
- 261 Paranyushkin, D. (2018). *InfraNodus*. Nodus Labs. <https://infranodus.com/>
- 262 Solymos, P., Zawadzki, Z., Bengtsson, H., & Team, R. C. (2023). *Pbapply: Adding progress bar*
 263 *to '*apply' functions* (Version 1.7-2). [https://cran.rstudio.com/web/packages/pbapply/](https://cran.rstudio.com/web/packages/pbapply/index.html)
 264 [index.html](https://cran.rstudio.com/web/packages/pbapply/index.html)
- 265 Weischedel, R., Pradhan, S., Ramshaw, L., Kaufman, J., Franchini, M., El-Bachouti, M.,
 266 Xue, N., Palmer, M., Hwang, J. D., Bonial, C., Choi, J., Mansouri, A., Foster, M.,
 267 Hawwary, A., Marcus, M., Taylor, A., Greenberg, C., Hovy, E., Belvin, R., & Houston, A.
 268 (2012). *OntoNotes release 5.0 with OntoNotes DB tool v0.999 beta*. BBN Technologies.
 269 <https://catalog.ldc.upenn.edu/docs/LDC2013T19/OntoNotes-Release-5.0.pdf>
- 270 Wickham, H., François, R., Henry, L., Müller, K., Vaughan, D., Software, P., & PBC. (2023).
 271 *Dplyr: A grammar of data manipulation* (Version 1.1.4). [https://cran.r-project.org/web/](https://cran.r-project.org/web/packages/dplyr/index.html)
 272 [packages/dplyr/index.html](https://cran.r-project.org/web/packages/dplyr/index.html)
- 273 Wickham, H., Hester, J., Chang, W., Bryan, J., & RStudio. (2022). *Devtools: Tools to make*
 274 *developing r packages easier* (Version 2.4.5). [https://cran.r-project.org/web/packages/](https://cran.r-project.org/web/packages/devtools/index.html)
 275 [devtools/index.html](https://cran.r-project.org/web/packages/devtools/index.html)
- 276 Wickham, H., Vaughan, D., Girlich, M., Ushey, K., Software, P., & PBC. (2024). *Tidyr: Tidy*
 277 *messy data* (Version 1.3.1). <https://cran.r-project.org/web/packages/tidyr/index.html>